

Assignment 1 - 2D Filter in OMP

Sanyam Gupta (sanyamgu)
10/17/2021

Overview

The task here is to implement a 2D filter using OMP so that the calculation can be spread across multiple cores. Analysis of the algorithm is also provided by repeating the calculation for 3 different matrix sizes. The speedup for [2, 4, 6, 8, 10] cores has been analysed. The reported data has been collected from JobID - **8130740**. The slurm.sh and 8130740.stdout files are included with the submission.

Algorithm for applying 2D filter

The overall step is broken down into 2 parallel computations.

1. Apply the filter in parallel and store the result in an auxiliary array.

```
std::vector<float> aDash(n*m);  
#pragma omp parallel for shared(K,A,n,m)  
for (int index=0; index<n*m; index++){  
    // if corner indices, value is original  
    if ((index/m == 0) | (index/m == n-1) | (index%m == 0) || (index%m == n-1)){  
        aDash[index] = A[index];  
    }  
    else {  
        aDash[index] = transformation(index, n, m, K, A);  
    }  
}
```

Figure 1 - Applying filter and store result in aDash

2. Copy the values from the auxiliary array to the original array.

```
#pragma omp parallel for shared(A, aDash)  
for (int index=0; index<n*m; index++){  
    A[index] = aDash[index];  
}
```

Figure 2 - Updating original array(A) with Adash

Ascertaining performance

The performance on 3 different arrays sizes, across 6 cores has been analysed. The base case i.e. without OMP support has also been reported. For each measurement, the experiment was repeated 5 times.

After discarding the maximum and minimum value, an average of **runtime** across 3 readings is shown below. The **speedup** data is also reported.

	Times (s)			CORES			
		1 (without omp)	2	4	6	8	10
	10000 * 20000	5.395	3.125	1.911	1.479	1.24	1.069
SIZE (n*m)	10000 * 40000	10.784	6.057	3.547	2.792	2.268	2.012
	10000 * 60000	16.184	9.033	5.092	4.014	3.368	2.885

Figure 3 - execution times

	Speedups			CORES			
		1 (without omp)	2	4	6	8	10
	10000 * 20000	1	1.726	2.823	3.648	4.351	5.047
SIZE (n*m)	10000 * 40000	1	1.78	3.04	3.862	4.755	5.36
	10000 * 60000	1	1.792	3.178	4.032	4.805	5.61

Figure 4 - Speedup values




Figure 5 - Plot for Speedup Values

Analysis of the Data

- **Strong Scaling** - For the same problem size (n) as the number of processing(p) elements decreases the runtime decreases. Hence the algorithm is strongly scalable.
- **Weak Scaling** - The runtime is also preserved across n/p values. Also a strongly scalable algorithm is also weakly scalable.